Math and Theory: collection overview

Dev D. Goyal

29 October 2026

Autonomous AI research disclosure. This collection reports fully autonomous agentic research. Each paper records its supported contribution history. Dev D. Goyal is responsible for publication. No external peer review is claimed.

Abstract

The catalogue contains thirty manuscripts and research notes across 29 result families in mathematics, theoretical computer science, and mathematical physics. Each paper states its formal verification scope. The thermal proof covers finite spectral statements, without a formal stochastic long-time identification theorem. The Gyires preprint proves inequalities for structured matrices and develops related polynomial bounds. Its principal order-five and order-six claims, the all-order permutation-polynomial and four-block theorems, and the exceptional-vector theorem have complete Lean proofs. The results comprise counterexamples, exact special cases, and sharp sequence bounds. The collection makes no claim of first priority. The papers have not been peer reviewed.

Contents

1 How to use the collection **2**

2 Mathematics **2**
2.1 Matrix theory . . . 2
2.2 Graph theory and graph polynomials . . . 3
2.3 Inequalities and real analysis . . . 4
2.4 Integer sequences and combinatorics . . . 4
2.5 Discrete analysis, probability and entropy . . . 5
2.6 Matroid theory . . . 5
2.7 Optimization . . . 5

3 Theoretical computer science **6**
3.1 Sampling algorithms and information theory . . . 6
3.2 Boolean functions and decision trees . . . 6
3.3 Fair division . . . 7
3.4 Proof complexity . . . 7

4 Mathematical physics **7**
4.1 Stochastic thermodynamics . . . 7
4.2 Quantum information . . . 8

5 Selected preprint packages **9**

6 Verification and authorship **9**

1 How to use the collection

The manuscript catalogue links each paper to its formal statements and independent computations. The verification guide records the commands and their results. Seven families share a Lake project pinned to Lean 4.35.0. Its 29 selected Comparator statements include equality, sharpness, and auxiliary positivity claims. The decision-tree and thermal packages use separate Lean 4.19.0 environments and a pinned Mathlib revision. The decision-tree package has 29 axiom reports of its own. This is a separate count from the shared Comparator selection.

2 Mathematics

2.1 Matrix theory

2.1.1 Real positive-definite counterexamples to extended q-permanent monotonicity

For every real parameter $q > 1$, some real symmetric positive-definite matrix has a negative q-permanent derivative at q. An explicit matrix of order five decreases from $q = 14$ to $q = 15$. The construction also gives correlation matrices with no zero entries and integer-entry examples. The dimension can depend on q. The real-symmetric question on $[-1, 1]$ remains outside the result.

Scope: Counterexample and all-parameter extension.

Manuscript Formal statement and prior work.

2.1.2 Gyires's permanent inequality for structured matrices

Every nonnegative real circulant matrix A of order five satisfies $24\,\mathrm{per}(AA^T) \le 625\,\mathrm{per}(A)^2$. Equality holds exactly for constant matrices. This proves the order-five circulant case of Gyires's ratio inequality. A bound in sorted-gap coordinates has sharp coefficient 484. The conjecture for arbitrary doubly stochastic matrices remains open.

Scope: Structured matrix inequalities and related polynomial bounds. The principal order-five and order-six claims, the all-order permutation-polynomial and four-block theorems, and the exceptional-vector theorem have Lean proofs.

Manuscript Formal statement and prior work.

Broader arguments in the same preprint

The manuscript also proves results for circulants of orders six and seven, with sharp sorted-gap margins. It treats $aI + bQ + cQ^2$ in every dimension, where Q is any permutation matrix and the coefficients are nonnegative. Further results cover four-block doubly stochastic matrices with unequal block sizes and boundary entries, matrices with one exceptional column or row, and every rank-at-most-two doubly stochastic matrix through order fifteen.

The supporting sections prove structural, scalar, deletion, balancing, local polynomial, and configuration-model results. Exact counterexamples identify the limits of auxiliary assertions. The complete support record gives the source map and exact reproduction commands.

The all-order four-block and exceptional-vector theorems have complete Lean proofs. Their inequalities and equality cases are among the 29 selected formal statements. The rank-two theorem has a complete Lean proof for orders 2 through 8. The general order-seven circulant theorem and rank-two orders 9 through 15 remain incomplete in Lean. Separate Lean packages prove the principal order-six and all-order permutation-polynomial theorems. The all-order rank-two problem and the unrestricted Gyires conjecture remain unproved here.

2.1.3 A polygon enclosure fails for sums of normal matrices

A polygon enclosure fails for a specified pair of normal matrices and its complete transposition graph. Actual normal matrices, determinants, straight-segment exterior and component exclusion are formalized. Correspondence with an external source conjecture is conditional.

Manuscript Reproduction and source context.

2.1.4 An unbounded obstruction to projection AM–GM in Loewner order

Three rank-one orthogonal projections in dimension three violate the symmetrized-product Loewner AM–GM bound. No finite real multiplier repairs it, even with a positive-definite mean. The examples do not satisfy the additional nonsingular block structure of the motivating coordinate-descent application.

Formal scope: Lean verifies a rational three-projection matrix obstruction and a parameterized obstruction for every real multiplier in dimension three; rank-one, positive-definite mean, rational all-multiplier choices, and the standard minimal-factor observation have complete conventional proofs.

Manuscript Proof and reproduction.

2.2 Graph theory and graph polynomials

2.2.1 An entropy increase for a structurally disassortative DomiRank graph

A connected graph with 14 vertices and 32 edges has negative degree–log-degree edge covariance. All resolvent scores are positive on $[251/1000, 63/250]$ inside the stable domain. Its entropy nevertheless increases by more than $1/10000$ in natural-logarithm units. This refutes the source's assertion for every positive-score interval.

Scope: Counterexample to an interval monotonicity conjecture.

Manuscript Formal statement and prior work.

2.2.2 Domination sequences of powers of paths and cycles

The domination counts of every natural power of a path or cycle form an ultra-log-concave sequence. Dividing each count by its binomial coefficient gives the normalized log-concavity inequality. The proof includes all natural graph orders and powers. It does not assert the same property for arbitrary graphs or real-rootedness of their domination polynomials.

Manuscript Formal proof and reproduction.

2.2.3 Sharp escape of complex domination roots in square flowers

Sharp complex-root escape for square flowers, with complete flower and book real-root classifications and multiplicities. All numbered manuscript conclusions are formalized in seventy modules with 258 endpoint audits. Historical novelty and exact external-source correspondence remain provisional.

Manuscript Reproduction and source context.

2.2.4 Leaf deletion can move the unique mode of a tree domination polynomial by two

Deleting a leaf can move the unique mode of an explicit tree domination polynomial by two. The fixed tree, deletion, actual domination counts and unique-mode movement are formalized. External question identity and historical priority remain provisional.

Manuscript Reproduction and source context.

2.3 Inequalities and real analysis

2.3.1 Diagonal concavity does not imply concavity for sums of elementary symmetric polynomials

For $M > 0$, let $f_M = 1 + M\sigma_1 + \sigma_2 + M\sigma_3 + \sigma_4$ in four variables, where σ_j is the elementary symmetric polynomial of degree j. The fourth root $f_M^{1/4}$ is concave on the positive diagonal exactly when $M \geq 1$. It is concave on the full positive orthant exactly when $M = 1$. Thus every $M > 1$ disproves Lachaume's criterion with strictly positive coefficients.

Scope: Counterexample and complete parameter classification.

Manuscript Formal statement and prior work.

2.4 Integer sequences and combinatorics

2.4.1 Alternating Motzkin sums

The principal root bound and bounded-forcing companion form one sequence family.

Sharp root log-concavity of alternating Motzkin sums Let $T_n = \sum_{j=0}^{n}(-1)^{n-j}M_j$, where M_j is the Motzkin number. The sequence $T_n^{1/n}$ is strictly log-concave on its tail starting at five. The inequality reverses at center five, so no earlier tail is log-concave. This strengthens Zhao's proposed tail starting at nine. The source recurrence and finite sum are linked in Lean.

Scope: Conjecture resolution with an optimal threshold.

Manuscript Formal statement and prior work.

Sharp log-concavity for alternating Motzkin sums A robust bounded-forcing recurrence theorem and corrected scaled corollary, with applications to alternating Motzkin-type sums. The arbitrary-real recurrence theorem and scaled corollary, including zero forcing, are formalized. Concrete sequences, uniform-family applications, root transfer and threshold classification retain conventional proofs.

Manuscript Reproduction and source context.

2.4.2 Counterexamples to a late-iterate ordering conjecture for generalized Hofstadter functions

For every $k \geq 4$, some input n satisfies $F_k^{\circ(2k+1)}(n) > F_{k+1}^{\circ(2k+2)}(n)$. Here the superscript denotes repeated composition. The proof constructs inverse iterates with exact gap $k-3$. The witness $F_4^{\circ 9}(132) = 8 > 7 = F_5^{\circ 10}(132)$ gives the first instance. The result concerns infinitely many parameters. It does not assert infinitely many violating inputs for each fixed parameter.

Scope: Counterexamples for every parameter k≥4.

Manuscript Formal statement and prior work.

2.4.3 A positive constant-recursive sequence with nonquasipolynomial power ranks

Exact minimal recurrences for every power of a rank-thirteen sequence, nonquasipolynomial ranks, and infinitely many power-rank profiles at each rank at least three. The main and arbitrary-rank companion theorems are formalized. Generating-function and polynomial corollaries retain conventional proofs.

Manuscript Reproduction and source context.

2.5 Discrete analysis, probability and entropy

2.5.1 The sharp Fourier bound at density one eighth

For every indicator of mean 1/8 on a finite uniform Boolean cube, the level-one Fourier weight is at most 3/64. A codimension-three subcube attains this bound in every dimension at least three. The proof also gives the equivalent sharp average Hamming-distance bound. It uses exact rational polynomial certificates and a dimension-independent moment identity.
Manuscript Formal proof and reproduction.

2.5.2 A quartic entropy deficit for random rankings

Let P be a probability distribution on the complete rankings of n labels. If every pairwise comparison has probability at distance at least $\varepsilon > 0$ from 1/2, then

$$H(P) \le (1 - \varepsilon^4/144)\log(n!).$$

The formal theorem includes every finite size, including zero and one. It assumes neither independent comparisons nor transitive preferred directions. It covers both ranking lists and position-valued permutations.

Manuscript, Lean proof, and reproduction commands.

2.5.3 Sharp universal Rényi threshold for iid sums on finite Sidon sets

A sharp universal Renyi-order threshold for iid sums supported on finite strong-B2 sets. The finite-support iid theorem, zero probabilities and critical endpoint are formalized. Comparison with unverified source definitions is conditional.

Manuscript Reproduction and source context.

2.6 Matroid theory

2.6.1 The asymmetric Merino–Welsh region

The theorem characterizes every nonnegative pair (x, y) for which

$$T_M(x, 0)T_M(0, y) \ge T_M(1, 1)^2$$

holds for all finite matroids without loops or coloops. The criterion uses two explicit scalar inequalities for every integer $k \ge 2$. The proof includes empty ground sets and finite ground sets in arbitrary ambient types. The package fetches Liu's pinned proof library from its public repository.
Manuscript Formal proof and reproduction.

2.7 Optimization

2.7.1 A rational five-cycle for two-block ADMM with multiplier step below two

An exact period-five orbit for multiplier-only two-block ADMM at step 199/100. Seven modules and eleven endpoint audits cover the orbit, minimizers and optimizer. Analytic regularity of the finite strongly convex extension retains a conventional proof.

Manuscript Reproduction and source context.

2.7.2 The exact last-gradient bound for constant-step gradient descent in the strongly convex case

The exact final-gradient guarantee holds on every real Hilbert space for $L > 0$, $0 \le \mu \le L$, normalized step $0 \le h \le 2$ and every natural horizon. Actual one-dimensional Huber and quadratic functions attain its two branches.

Formal scope: The actual-function upper bound and normalized actual-function attainment are formalized, including every horizon and closed parameter boundary. Sharpness for arbitrary $L > 0$ and radius $R > 0$ is proved by rescaling in the manuscript, without a separate Lean wrapper.

Manuscript Proof and reproduction.

3 Theoretical computer science

3.1 Sampling algorithms and information theory

3.1.1 Counterexamples to a best-of-n win-rate bound at every budget

For every $n \ge 2$ and $\delta > 0$, a finite positive reference distribution with largest atom at most δ violates the proposed best-of-n divergence–win-rate curve. A four-response best-of-two example gives exact probabilities and a strict logarithmic certificate. The all-budget proof uses finite interval bounds and gives a quantitative positive gap. It does not determine the sharp replacement curve.

Scope: Disproof for every budget $n \ge 2$ and every positive atom cap.

Manuscript Formal statement and prior work.

3.2 Boolean functions and decision trees

3.2.1 Logarithmic influence bounds and variance counterexamples for Boolean decision-tree covariance

Let T be a finite Boolean decision tree under uniform independent bits. No coordinate repeats on a root-to-leaf path. Let $k_i(T)$ count all query nodes labelled i, and let $I_i(f)$ be the probability that flipping coordinate i changes the output. The weighted sum of absolute child covariances satisfies

$$\sum_v 2^{-d(v)} \, |\mathrm{Cov}(f_{v,+}, f_{v,-})| \le \sum_i I_i(f) \log_2 \max\{1, k_i(T)\}.$$

Here $d(v)$ is the node depth, and the two child functions use the same input cube. The read number counts nodes in the entire tree, not queries along one path.

For a read-k tree, the bound is $I(f) \log_2 k$. The dummy-prefix construction of Wan, Wright, and Wu attains equality when k is a power of two. The logarithmic potential is due to Shalev. A four-variable counterexample has covariance $5/4$ and variance $3/4$. A nine-variable counterexample has covariance $4863/4096$ and variance $3007/4096$, with every query relevant. Both refute the coefficient-one logarithmic variance proposal for a given tree.

Scope: Covariance inequalities, an equality family, and exact finite witnesses. No Fourier-entropy theorem is formalized.

Manuscript Prior work Formal proof and reproduction.

3.3 Fair division

3.3.1 Ordered-instance allocation transfer with exact extra-copy preservation

A set-valued ordered-instance transfer preserving utility improvement, the exact extra-copy count, per-good caps and coverage. The transfer theorem is formalized. Existence of a fair allocation before transfer is outside the result; external source comparison is provisional.

Manuscript Reproduction and source context.

3.4 Proof complexity

3.4.1 Unbounded coalesced-mass certificates for resolution trees

A coalesced dyadic clause-mass functional has exact supremum n+1 on the stated unrestricted resolution-tree class and is unbounded across dimensions. The internally defined resolution-tree functional and proof class are formalized. No unconditional identification with an external conjecture or a regular proof class is claimed.

Manuscript Reproduction and source context.

4 Mathematical physics

4.1 Stochastic thermodynamics

4.1.1 A three-state thermal counterexample to a heat-current kurtosis conjecture

A three-state rate model with two positive-temperature baths satisfies local detailed balance. Every total off-diagonal rate is positive. The tilted generator records heat transferred to the cold bath. Its dominant eigenvalue branch is four times continuously differentiable near zero and has a larger real part than every distinct complex eigenvalue. The exact derivative ratio is

$$\frac{C_4}{C_2} = -\frac{3470947982622836}{176627429595846691} < 0,$$

where C_n is the nth derivative at zero. This refutes Eq. (19) of Ptaszyński, *Physical Review E* **106**, 024119 (2022), in the paper's finite thermal spectral formulation. The third derivative is positive, so the witness does not refute the separate skewness conjecture.

Scope: Finite rate matrices, local detailed balance, complex eigenvalues, and derivatives of a local dominant branch. The stochastic long-time theorem that identifies the branch with a cumulant generating function is not formalized. No minimal-state or empirical validation claim is made.

Manuscript Prior work Formal proof and reproduction.

4.1.2 A sharp universal skewness bound for unicyclic currents

The net marked-edge current on every finite bidirectional Markov ring with at least three states, positive rates and unequal forward and reverse rate products satisfies $-1/16 \leq \mathrm{C}_3/\mathrm{C}_1 \leq 1$. The two-channel two-state case has the stronger interval [1/4,1]. This ratio is not standardized skewness.

Formal scope: The interval for actual tilted-matrix derivatives and a local branch is formalized. Identification with long-time stochastic cumulants and sharpness have conventional proofs. Lean does not construct a continuous-time sample-path probability space or a separate optimizing-sequence theorem.

Manuscript Proof and reproduction.

4.2 Quantum information

4.2.1 Universal entanglement embezzlement

A target-independent catalyst of local dimension 2^{4L} gives root infidelity at most $125m/L^2$ for every Schmidt probability vector on m entries and $L \geq 1$. Zero entries are permitted. The proof constructs local basis permutations and gives their finite complex Hilbert-space interpretation. It does not separately formalize arbitrary changes to an external Schmidt basis, gate efficiency, or catalyst reuse.
Manuscript Formal proof and reproduction.

4.2.2 The four-qubit Higuchi–Sudbery entropy bound

For every normalized pure state of four qubits, the average von Neumann entropy of the three balanced two-qubit cuts is at most $\log 2 + \frac{1}{2}\log 3$. The Higuchi–Sudbery state attains this value. The formal proof uses all sixteen complex amplitudes, the actual reduced matrices and their eigenvalues, an exact sum-of-squares identity, and an explicit attaining state. It does not classify all maximizers or cover mixed states.

Manuscript, Lean proof, and reproduction commands.

4.2.3 A pure-state matched-basis information inequality violation in dimension five

A dimension-five pure-state matched-basis information inequality violation, with a conventional all-prime companion. The dimension-five identical-local-basis theorem is formalized. The prime-family extension is conventional; arbitrary or independently optimized basis choices are outside the result.

Manuscript Reproduction and source context.

4.2.4 Exact three-setting correlation bounds at fixed steering centre

The exact least upper bound of the optimized three-setting squared correlation value is determined for every two-qubit steering-centre length. The threshold for a value greater than one is $c < 2/3$. Attainment, nonattainment and the singular marginal boundary are included.
Formal scope: Actual two-qubit density matrices, canonical filters, Pauli correlations, the exact least upper bound, attainment and nonattainment, and fully entangled fraction are formalized. The steering-ellipsoid and optimized CJWR measurement interpretations use cited external correspondences; general steerability is outside the result.

Manuscript Proof and reproduction.

4.2.5 The three-state root-fidelity entropy bound

The entropy of the weighted three-by-three root-fidelity matrix bounds the Holevo quantity of every three-state ensemble in arbitrary finite dimension. The result includes mixed and singular states and zero probabilities.
Formal scope: The unrestricted three-state ensemble theorem is formalized for actual positive semidefinite complex matrices of trace one, nonnegative weights summing to one, spectral entropy and nuclear-norm root fidelity. Credited QICLean foundations and prior-work adapters retain their attribution.

Manuscript Proof and reproduction.

5 Selected preprint packages

The repository includes a ZIP containing nine selected preprints: Higuchi–Sudbery, universal embezzlement, asymmetric Tutte, power-rank, ADMM, steering bounds, the strong-convexity gradient bound, unicyclic skewness and three-state root fidelity. Each has a first-page autonomous AI disclosure, source archive, metadata and upload instructions. No external upload or DOI registration is claimed.

6 Verification and authorship

Lean checks the proofs against their formal statements. In the shared Lean 4.35.0 library, separate Challenge modules state the intended claims. Comparator checks these against their proved Solutions. The verification record identifies the completed runs of Lean, NanoDa, and con-ron for those seven families. The standalone study packages have their own pinned Lean environments and stated proof boundaries. Source-integrity checks are distinguished from completed compilation records. Their audit logs state the permitted axioms. Exact arithmetic programs check the accompanying finite certificates. The shared Comparator results do not cover those packages. Independent exact and numerical programs test finite certificates and examples. These finite computations do not replace the general proofs.

This work is fully autonomous agentic research. The original collection records GPT 6 with Pro Reasoning in ChatGPT for initial research, GPT 6 Astra in Codex for formalizations and corrections, and the author's stated proof check. These historical statements are not inferred for later uploads. The later packages state their own supported contributions, including actual GPT 6.1 Sol formalization and verification work. Dev D. Goyal is responsible for publication. No expert human validation or external peer review is claimed for the later work. The authorship disclosure identifies these tools and their roles. These preprints have not been peer reviewed.